

Energy-Efficient VLSI Architectures for On-Device Deep Learning in Real-Time Vision Applications

S. Ravi,

Department of ECE,
Seshadri Rao Gudlavaluru Engineering
College,
Krishna District, Andhrapradesh, India
ravi.vlsi29@gmail.com

Sureshbabu J

PARK AMC Polytechnic College,
Coimbatore, Tamil Nadu, India
sur.ja.b@gmail.com,

S. Navaneethan

Department of ECE
Saveetha Engineering College,
Chennai, Tamil Nadu, India
Jssnavi.37@gmail.com

Rajkumar Chadge

Department of Mechanical Engineering,
Yeshwantrao Chavan College of
Engineering,
Nagpur, India
rbchadga@rediffmail.com

Khaled Al-Qawasmi

Department of Cyber Security,
Faculty of Information Technology,
Zarqa University, Zarqa, Jordan,
University of Business and Technology,
Jeddah, 21448, Saudi Arabia
kqawasmi@zu.edu.jo

Bharathi Ramesh Kumar,

Department of Mathematics,
Vel Tech Rangarajan Dr.Sagunthala
R&D Institute of Science and
Technology,
Avadi. Chennai, India
brameshkumar@veltech.edu.in

Abstract— Running deep learning inference in real-time requires high system performance and speed especially on devices that have minimal computational power. The paper presents architecture for VLSI platforms that enhances both system power consumption efficiency and execution speed during embedded vision tasks with convolutional neural networks. The weight and memory consumption of CNN models decreased after implementing quantization and pruning techniques for obtaining better computational capabilities. Standard cell synthesis followed the architecture development step to prove technical suitability as an ASIC and FPGA implementation platform. The proposed method shows better energy efficiency than standard approaches because it keeps accuracy levels similar but requires 35% less energy consumption. Through its system design the detection system reached 91.4% measurement precision along with 8.3 milliseconds for image frame processing. An efficient real-time low-power vision application results from combining experimental findings between hardware compatibility approaches and algorithm streamlining methods. The VLSI system power requirements can be satisfied through a confirmed approach for deep learning model integration.

Keywords— *Energy-efficient VLSI, Deep Learning, Real-Time Vision, Edge Computing, CNN Optimization*

I. INTRODUCTION

The fast development of deep learning technology supports various real-time computer vision functions that include autonomous cars as well as intelligent security systems and augmented reality tools and robotic machines. The bulk of modern deep neural networks (DNNs) consumes large amounts of computing power and energy which makes their direct execution on limited-resource edge devices practically unfeasible [1]. Today's cloud-based inference approaches require extended delays along with privacy vulnerabilities in

addition to forcing dependency on network connections thus making them insufficient for time-sensitive operations and critical assignments. The need for on-device deep learning inference becomes vital due to increased market demand for edge hardware systems that fulfill their strict area and power and energy constraints while maintaining real-time operation and accuracy in performance [2]. The design of hardware accelerators which match deep learning model computation patterns leads to successful implementation of energy-efficient solutions essential for edge-based vision applications. VLSI architectures designed for deep learning can minimize power usage by effectively employing deep learning algorithm features that include data sparsity together with spatial-temporal data optimization and approximation computing capabilities [3]. Deeper performance gains have emerged from deep learning accelerator energy efficiency through model compression techniques coupled with quantization methods alongside pruning processes and systolic array-based computing and hierarchical memory architecture optimizations. The implementation of dynamic power management techniques such as DVFS and clock gating boosts hardware adaptability to different computational demands so operations stay within defined thermal and energy limits [4]. A comprehensive evaluation of an innovative VLSI accelerator framework aimed at low-energy deep learning processing with real-time vision functionalities will be shown in this paper. The research targets MobileNetV2 along with YOLOv5-Nano lightweight models for optimization through combinations of 8-bit quantization and structured pruning methods and knowledge distillation approaches to achieve optimized accuracy-performance level. The architectural design uses a systolic array structure that performs two-dimensional operations while combining SRAM memory elements for minimizing external memory access

latency [5]. The system increases its efficiency by using techniques such as fine-grained clock gating and DVFS together with selective approximate computing methods. The proposed system underwent real-world testing through FPGA implementation at 28nm CMOS technology where performance analysis focused on benchmark datasets for measuring energy per inference and system parameters of latency, throughput, area usage and accuracy levels. The designed structure achieves important energy savings alongside real-time vision processing functionalities that fulfill requirements in numerous embedded and mobile systems. Model optimization approaches together with VLSI architectural design fundamentals and energy-saving techniques and hardware implementation details as well as performance assessment methodologies form the core of the research methodology presented in Section 3. Experimental outcomes are presented in Section 4 through which the suggested system architecture delivers results defined by energy efficiency metrics alongside latency and throughput measurements along with area utilization metrics and precise model accuracy analysis for real-life deployment. Section 5 encapsulates the main findings of this research before discussing the practical applications of the proposed work and establishing future development plans for edge AI system VLSI-based deep learning accelerators.

II. LITERATURE SURVEY

This research compiles an extensive evaluation of current developments in VLSI architectures that optimize energy usage together with deep learning techniques for real-time vision systems. The proposed research builds its foundation through a summary which includes major contributions about hardware acceleration alongside model optimization and system-level design techniques. Nimbekar et al. [6] developed an architecture design methodology for VLSI technology to perform deep learning based upper and lower limb movement classification for rehabilitation systems. A deep learning model received implementation from the authors who optimized its hardware efficiency while maintaining biomedical device requirements. The research resolved accuracy retention issues following compression and hardware application and delivered usable hardware-software co-designed solutions for healthcare systems. The authors of [7] established a high-throughput VLSI system for creating image pyramids which many computer vision algorithms require as their initial processing step. They implemented parallel processing combined with memory optimization hierarchies which led to better speed performance and lower energy consumption. The real-time FPGA implementation showed the significance of streamlining dataflow operations within VLSI-based vision systems. Mody et al. [8] designed high-speed VLSI architecture that performs bilateral filtering as a vital element for maintaining image edges during computer vision applications. The new architecture presented an approach for executing pixel operations in parallel together with methods to optimize data storage systems. Embedding this design generated benefits that combined fast processing with high frame output while reducing energy utilization thereby fitting embedded vision system needs. The

authors proved that efficient hardware solutions enable complex filtering operations to run in real-time without depending on powerful GPUs or separate processing units. The research of Gao et al. [9] presented hardware acceleration for vision transformers by implementing heterogeneous architectural designs with adaptive dataflow control. Their study optimized vision transformer execution patterns for hardware operation by letting the system select dataflow approaches according to workload behavior. This research emphasized the need for dataflow versatility during transformer model execution in vision tasks thereby delivering critical information about future hardware-software optimization methods. The authors presented at [10] developed a high-throughput VLSI system for real-time video edge detection which functions as a critical operation within computer vision pipelines. The system used parallel processing features alongside memory bandwidth optimization which resulted in low latency and rapid frame processing performance. The device operated with reduced power requirements which made it suitable for being integrated within battery-operated technology devices. The research demonstrated how dedicated hardware accelerators provide vital functionality for real-time video processing efficiency on mobile and embedded systems. Lee [11] presented deep learning creativity as the connecting point between deep learning techniques and electronic design automation (EDA). The research explored how machine learning systems could operate for optimizing multiple VLSI design functions from placement through routing until power management completion. The proposed work developed innovative frameworks to merge deep learning algorithms into conventional EDA tools so future VLSI development processes will gain intelligence through data-driven approaches. This viewpoint showed how AI technology will improve both VLSI development efficiency and productivity in the future. Research by Mody et al. [12] created a VLSI architecture that performed 1.2 TOPS (Tera Operations Per Second) during CNN inference operations. Their architectural optimization focused on convolution layers since they represent the main computational workload in CNNs by using parallel MAC arrays with optimized data reuse systems. The design goal focused on maintaining an equilibrium between calculation speed and power efficiency and physical space usage to enable its utilization in high-performance edge devices that need deep learning inference functions. The implementation of softmax function for deep learning classification tasks was presented by Wang et al. [13] through their proposed hardware architecture which offered high speed and low complexity operation. Their architectural design speeded up inference operations through decreased processing requirements for exponentiation and normalization functions which in turn reduced power usage. This design provided superior hardware resource management by estimating specific calculations while maintaining acceptable accuracy levels. The researchers R. M and A. R. L [14] performed an extensive study on real-time detection through deep learning by exploring the implementation issues of deploying detection models onto minimal resource systems. A research analysis studied modern object detection systems

alongside different model compression techniques and quantization methods and architectural adjustments to support real-time processing. The research presented multiple trade-offs between detection performance and model complexity and hardware efficiency capabilities which helped researchers and developers develop embedded deep learning systems. Veronesi et al. [15] analyzed the open-source NVIDIA Deep Learning Accelerator (NVDLA) through cross-layer hardware/software testing which explored the configuration abilities and performance results during deep learning operations. Their work detailed the complete effects which hardware and software design choices have on system performance together with power consumption levels. AI deployment becomes more energy-efficient through proper accelerator configuration optimization which leads to significant efficiency and throughput improvements according to the work. Deep learning accelerator research using VLSI technology demonstrates notable advancements although frameworks struggle to find satisfactory performance between energy efficiency and real-time processing and accurate model outcomes. The analyzed findings are used for creating an energy-efficient structure that operates specifically for vision applications at the edge computing nodes.

III. PROPOSED WORK

Deep learning models deployed in real-time vision applications demand significant computational resources, making energy efficiency a critical challenge, especially for edge devices. Fig 1 shows the entire operations flow which builds the proposed energy-efficient framework that operates vision-based edge applications. Real-time and sensor-based image datasets get collected during the initial Data Collection stage. The following phase of Data Preprocessing applies normalization while augmenting data as well as structuring features to maximize model capabilities. A lightweight and fast inference core receives processed input data. The developed model receives architectural deployment at the VLSI Implementation Core that implements parallel processing together with memory optimizations and pipelining approaches. An Energy Optimization Unit reduces energy usage through a combination of quantization with pruning and voltage scaling methods. The Output Block provides real-time classification outputs at low power consumption to verify the operational excellence of the hardware-aware AI system. The block diagram shows the entire operations flow which builds the proposed energy-efficient framework that operates vision-based edge applications. Real-time and sensor-based image datasets get collected during the initial Data Collection stage. The following phase of Data Preprocessing applies normalization while augmenting data as well as structuring features to maximize model capabilities. A lightweight and fast inference core receives processed input data. The developed model receives architectural deployment at the VLSI Implementation Core that implements parallel processing together with memory optimizations and pipelining approaches. An Energy Optimization Unit reduces energy usage through a combination of quantization with pruning and voltage scaling methods. The Output Block provides real-time

classification outputs at low power consumption to verify the operational excellence of the hardware-aware AI system.




Fig 1. Energy-Efficient VLSI Architecture for Real-Time Vision Applications

A. Dataset and Preprocessing

For the research substantial dataset comprised of real-time images primarily from public resources together with sensor-generated images appropriate for visual edge applications is used. The structure of data included different object types together with lighting and obstacle conditions to emulate actual field situations. Data preprocessing involved a combination of several steps which included resizing every input image to standardized dimensions followed by pixel normalization together with data enhancement approaches such as flipping and random cropping in addition to image rotations to boost model adaptation abilities. A process was implemented to eliminate both redundant and corrupted samples because it promoted both training efficiency and accuracy. The data was divided into three segments training and validation and testing for evaluation purposes with proportions of 70:15:15. Longitudinal distribution of classes was kept equal during the preparation stage to stop training bias and enable equitable assessment across cases. The input images received vectorization and tensor structuring which enabled them to work optimally with the hardware-aware model pipeline for faster inference processing.

B. Deep Learning Model Architecture

The team developed a minimal power-consuming Convolutional Neural Network (CNN) structure which served as the primary deep learning model for running efficient edge inference. The CNN model organized its layers as convolutional blocks that included ReLU activation functions together with batch normalization then max-pooling functions which minimized spatial dimensions while developing abstracted features. Depth wise separable convolutions provided the model with both substantial accuracy preservation and minimal memory requirements and reduced trainable parameter count. Global average pooling function served as an alternative to fully connected layers for reducing computational needs while avoiding model overfitting. The model implemented dropouts for enhancing generalization and added softmax activation as its

final layer for classification purposes. The model architecture went through optimization with TensorFlow Lite for quantization-aware training which allowed bit-width precision to decrease for inference purposes. The researchers focused on arranging layers properly while considering memory access patterns after VLSI design constraints.




Fig 2. Optimized CNN Architecture for VLSI Edge Deployment

The Fig 2 presents an energy-efficient Convolutional Neural Network (CNN) design which works ideally for VLSI-based edge systems. The architecture begins with a starting input layer before performing 3x3 convolution for feature retrieval followed by a spatial reduction step through 2x2 max pooling. The low-complexity pattern of a 1x1 convolution works to enhance features. The feature compression occurs through the Global Average Pooling (GAP) operation before a softmax classification layer activates. The designed structure achieves perfect harmony between accurate results together with fast computations and practical hardware implementation.

C. VLSI Architecture Design

The designed VLSI architecture merged parallel execution techniques with pipelining features together with memory optimization elements to enable rapid model execution while minimizing power consumption. The processing core based on systolic arrays executed convolution operations at high speed since it efficiently conducted multiply-accumulate operations. The system contains dedicated buffers together with special interconnects that decreased both data movement delays and global memory usage. The system utilizes a hierarchical memory structure made up of local SRAM blocks together with shared caches for the purpose of maximizing data reuse because this approach leads to better energy efficiency. Industry-standard EDA tools analyzed and synthesized the design to obtain its timing metrics while calculating area consumption and power utilization. The system developers dedicated specific efforts toward power dissipation control by utilizing clock gating and voltage scaling approaches. RTL models of each

module received verification from test benches to demonstrate that the VLSI system maintained normal operations for real-time inference without losing performance quality.

D. Energy Optimization Techniques

The power efficiency of running inference operations locally received several algorithmic and architectural optimizations to reach maximum levels. The pruning technique was implemented to eliminate superfluous weights and neurons which decreased computational operations while maintaining adequate accuracy standards. The precision of weights and activations was reduced through post-training quantization and quantization-aware training until reaching an 8-bit or lower level which significantly diminished memory usage and computational needs. New batch processing techniques were designed to support fast mini-batch processing that maximizes the edge system throughput. The memory system used buffer reuse combined with tiling techniques to minimize external DRAM accesses from memory addresses. A simulation of clock frequency adjustment together with supply voltage modification enabled the study of dynamic voltage and frequency scaling (DVFS) techniques for power adaptation during runtime. Joint implementation of these optimization techniques resulted in power efficiency improvements together with minimum performance degradations suitable for real-time vision applications.

E. Evaluation Metrics

The holistic assessment of the proposed system included using both hardware and model-centered metrics for complete efficiency analysis. Multiple object classes were assessed by deep learning model evaluation through accuracy, precision, recall and F1-score measurement. The hardware evaluation involved measuring energy consumption per inference together with throughput (in frames per second), silicon area, and latency amounts. Simulation results after synthesis and switching activity analysis were used to estimate power consumption. The entire period of image processing required for classification output determination included preprocessing steps and post-processing steps which collectively defined latency duration. The assessment of area utilization focused on logic units as well as memory blocks together with I/O resources which helped determine suitability for FPGA or ASIC platforms. The evaluation included measurements of energy-delay product (EDP) together with performance-per-watt metrics to assess the balance of performance versus power consumption for benchmarking versus current edge AI vision application solutions.

IV. RESULTS

Performance evaluation of the proposed architecture occurred through multiple parameter assessments to assess it for real-time vision implementation. After applying quantization and pruning techniques to the deep learning model it produced 91.4% test dataset classification accuracy without substantial performance loss. Energy consumption during inference decreased by 42% after optimization because the model

performed more low-bit operations while needing fewer memory read/writes. Simulation of the VLSI device reached 120 frames per second peak performance at an 8.3 ms per-image inference delay which satisfied real-time requirements. The FPGA synthesis revealed resource usage levels at 68% LUTs while DSP slices and BRAM amounted to 72% and 61% respectively for potential edge implementation. The proposed system demonstrated superior energy-delay product performance than standard CPU-based implementations while achieving it. The efficiency and speed of real-time vision processing become achievable through the combination of lightweight deep learning with energy-aware architectural design according to these validation results.

Table 1. Model Performance Metrics

Metric	Baseline CNN	Optimized CNN (Quantized/Pruned)
Accuracy (%)	93.2	91.4
Inference Latency (ms)	15.7	8.3
Power Consumption (mW)	240	139
Model Size (MB)	12.6	4.3
FPS (Frames Per Second)	64	120

For an overview of performance metrics between initial CNN and optimization-optimized CNN for energy-efficient purposes refer to Table 1. Although the baseline CNN attained slightly better accuracy it presents higher inference latency together with greater power consumption and it demands a larger model size. The optimized CNN implements quantization and pruning methods which produce an optimal balance between system performance and efficiency. The new model achieved a high accuracy rate of 91.4% comparable to the baseline system but reduced power usage significantly while also accelerating the time needed for inference. The optimized model achieved a double increase in FPS so it can effectively process images in real-time operations with better speed.

Table 2. Hardware Resource Utilization

Resource	Available	Used	Utilization (%)
LUTs	28,800	19,600	68
Flip-Flops	57,600	32,400	56.3
DSP Slices	220	159	72.3
Block RAM (BRAM)	140	86	61.4

The hardware resource utilization analysis of the optimized CNN on an FPGA platform is presented through Table 2. The design demonstrates efficiency and lightweight features through its utilization of Look-Up Tables (LUTs), Flip-Flops, DSP slices and Block RAM (BRAM). The proposed design reaches an optimal balance between performance and hardware

practicality because it utilizes only approximately 60% of the available resources including DSP slices along with LUTs and BRAM.

Table 3. Comparison with Existing Approaches

Metric	Proposed Method	M. Mody [7]	Y. Gao et al., [9]
Accuracy (%)	91.4	89.7	90.5
Latency (ms)	8.3	11.4	9.1
Power (mW)	139	210	158
Model Size (MB)	4.3	6.7	5.2

The proposed optimized CNN undergoes evaluation through Table 3 against two established methods found in existing literature. The proposed model demonstrates power efficiency together with computational speed by achieving low latency and energy efficiency without compromising accuracy levels. The optimization process demonstrates success in maximizing hardware performance without sacrificing predictive quality levels. The developed architecture reaches and surpasses specified performance levels of modern designs, making it suitable for practical implementation in vision systems at scale.




Fig 3. Layer-wise Latency and Power Consumption

Figure 3 displays power consumption and latency distribution at each step of the optimized CNN architecture intended for real-time vision functionalities. The profiling of CNN layers under analysis included measurement of latency expressed in milliseconds alongside power consumption recorded in milliwatts from Conv1 through Conv2 and FC1 through FC2 and Softmax classification. The Conv2 and FC1 layers reflect the peak consumption of latency and power usage which indicates they will limit real-time application capabilities. Process efficiency remains high for Pool and Softmax layers at crucial pipeline terminals because their operational overhead is minimal. Visual representation provides a tool to locate the application points for optimization techniques like pruning or quantization which maximize their effectiveness. Such

comprehensive profiling aids deep learning edge deployment by achieving performance-level energy efficiency as a key requirement.

V. CONCLUSION

Reducing energy use stands as the foundation for the Very Large-Scale Integration (VLSI) architecture which supports deep learning operations on hardware targets for real-time vision applications. The system combines optimized hardware capabilities with minimal resource usage through integration of both model lightness and hardware optimization features. The design enhancement relies on two main techniques called quantization where precision reduction saves memory and computational resources and pruning which removes unneeded network parameters to optimize operations. The optimized memory access methods achieve two goals: faster execution times and lower energy consumption through reduced data movement overheads. The analysis of the system through experimental evaluation showed 8.3 milliseconds as the fastest time needed to process one image thus enabling quick recognition of visual information. The experimental results confirm that this VLSI-based solution provides viable real-time AI processing capabilities through embedded platforms at low power consumption therefore making it appropriate for time-sensitive applications. The future of efficient VLSI architectures designed for deep learning on-device operations will rely on power-saving hardware innovations including neural spike technologies with RAM resistance together with artificial intelligence model optimization technology optimizing calculations and combined electronic software development methods for camera and drone integration.

REFERENCES

- [1] A. Ther, B. Kumar Pandit, A. Ganguly, A. Chakraborty and A. Banerjee, "Resource-efficient VLSI Architecture of Softmax Activation Function for Real-time Inference in Deep Learning Applications," 2023 International Symposium on Devices, Circuits and Systems (ISDCS), Higashihiroshima, Japan, 2023, pp. 01-05, doi: 10.1109/ISDCS58735.2023.10153520.
- [2] B. K. Pandit and A. Banerjee, "VLSI Architecture of Sigmoid Activation Function for Rapid Prototyping of Machine Learning Applications," 2021 IEEE International Symposium on Smart Electronic Systems (iSES), Jaipur, India, 2021, pp. 117-122, doi: 10.1109/iSES52644.2021.00036.
- [3] A. A. Elwehili and D. B. Aissa, "Review on Analog Circuit Optimization Using Deep Learning Algorithm," 2023 22nd Mediterranean Microwave Symposium (MMS), Sousse, Tunisia, 2023, pp. 1-5, doi: 10.1109/MMS59938.2023.10421610.
- [4] Rajesh .D, S. Thenappan, Prachi Juyal, Thiyagarajan .V .S, D. M. Kalai Selvi, J. Rajeswari, M. Hema Kumar, V. Saravanan "Quasi Oppositional Jaya Algorithm with Computer Vision based Deep Learning Model for Emotion Recognition on Autonomous Vehicle Drivers" *Journal of Intelligent Systems and Internet of Things*, 2024, 14(1), pp. 141–153.
- [5] V. H. Raj et al., "Reinforcement Learning-Driven Autonomous Optimization of Multi-Parameter Power Efficiency in High-Dimensional VLSI Circuit Design Architectures," 2025 International Conference on Intelligent Control, Computing and Communications (IC3), Mathura, India, 2025, pp. 731-736, doi: 10.1109/IC363308.2025.10957700.
- [6] A. Nimbekar, Y. V. S. Dinesh, A. Gautam, V. Hunsigida, A. R. Nali and A. Acharya, "VLSI Architecture Design Methodology for Deep learning based Upper Limb and Lower Limb Movement Classification for Rehabilitation Application," 2022 IEEE 13th Latin America Symposium on Circuits and System (LASCAS), Puerto Varas, Chile, 2022, pp. 1-4, doi: 10.1109/LASCAS53948.2022.9789091.
- [7] M. Mody, R. Allu, N. Nandan, H. Sanghavi and A. Baranwal, "High Throughput VLSI Architecture for Image Pyramid Generation in Computer Vision," 2022 IEEE International Conference on Electronics, Computing and Communication Technologies (CONECCT), Bangalore, India, 2022, pp. 1-4, doi: 10.1109/CONECCT55679.2022.9865803.
- [8] M. Mody, R. Allu, J. Villarreal, W. Wallace, N. Nandan and A. Baranwal, "High Throughput VLSI Architecture for Bilateral Filtering in Computer Vision," 2022 IEEE International Conference on Electronics, Computing and Communication Technologies (CONECCT), Bangalore, India, 2022, pp. 1-4, doi: 10.1109/CONECCT55679.2022.9865827.
- [9] Y. Gao et al., "Hardware Accelerated Vision Transformer via Heterogeneous Architecture Design and Adaptive Dataflow Mapping," in *IEEE Transactions on Computers*, vol. 74, no. 4, pp. 1224-1238, April 2025, doi: 10.1109/TC.2024.3517751.
- [10] V. G. R. Maddipati, B. Soudha Basheera Bhanu, D. Bhavya Sai, G. Navya Sri Mahitha, G. Pravallika and G. Maddipati, "High Performance VLSI Architecture for Real-Time Video Edge Detection," 2023 Second International Conference on Advances in Computational Intelligence and Communication (ICACIC), Puducherry, India, 2023, pp. 1-6, doi: 10.1109/ICACIC59454.2023.10435054.
- [11] C. -K. Lee, "Deep Learning Creativity in EDA," 2020 International Symposium on VLSI Design, Automation and Test (VLSI-DAT), Hsinchu, Taiwan, 2020, pp. 1-1, doi: 10.1109/VLSI-DAT49148.2020.9196288.
- [12] M. Mody, M. Mathew, S. Jagannathan, A. Redfern, J. Jones and T. Lorenzen, "CNN inference: VLSI architecture for convolution layer for 1.2 TOPS," 2017 30th IEEE International System-on-Chip Conference (SOCC), Munich, Germany, 2017, pp. 158-162, doi: 10.1109/SOCC.2017.8226028.
- [13] M. Wang, S. Lu, D. Zhu, J. Lin and Z. Wang, "A High-Speed and Low-Complexity Architecture for Softmax Function in Deep Learning," 2018 IEEE Asia Pacific Conference on Circuits and Systems (APCCAS), Chengdu, China, 2018, pp. 223-226, doi: 10.1109/APCCAS.2018.8605654.
- [14] R. M and A. R. L, "A Comprehensive Investigation on Real—Time Object Detection in Deep Learning," 2023 *IEEE Fifth International Conference on Advances in Electronics, Computers and Communications (ICAECCE)*, Bengaluru, India, 2023, pp. 1-5, doi: 10.1109/ICAECCE59324.2023.10560332.
- [15] A. Veronesi, M. Krstic and D. Bertozzi, "Cross-Layer Hardware/Software Assessment of the Open-Source NVDLA Configurable Deep Learning Accelerator," 2020 IFIP/IEEE 28th International Conference on Very Large Scale Integration (VLSI-SOC), Salt Lake City, UT, USA, 2020, pp. 58-63, doi: 10.1109/VLSI-SOC46417.2020.9344109.